
MSiA 490 Lab 4

Fall 2023
Shuyang Wang

Agenda

- Announcement:
 - Part 1 of Assignment 1 is posted. Due on Oct. 19th at 10 pm (Thursday).
- Policy Gradient
- CartPole
- Pong
- Implementation

Policy Gradient

$$\nabla_{\theta} J(\pi_{\theta}) = \nabla_{\theta} \mathbb{E}_{\tau \sim \pi_{\theta}} [R(\tau)]$$

$$= \nabla_{\theta} \int_{\tau} P(\tau|\theta) R(\tau)$$

Expand expectation

$$= \int_{\tau} \nabla_{\theta} P(\tau|\theta) R(\tau)$$

Bring gradient under integral

$$= \int_{\tau} P(\tau|\theta) \nabla_{\theta} \log P(\tau|\theta) R(\tau)$$

Log-derivative trick

$$= \mathbb{E}_{\tau \sim \pi_{\theta}} [\nabla_{\theta} \log P(\tau|\theta) R(\tau)]$$



Return to expectation form

$$\therefore \nabla_{\theta} J(\pi_{\theta}) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t|s_t) R(\tau) \right]$$

Expression for grad-log-prob

CartPole & Pong

- See demo in the notebook.



Training a policy network

Policy(S_t) is a distribution over actions. You can build the distribution using softmax function with logits.

A Neural Network as the Policy Network: Input: State; Output: Logits.

Pseudocode:

1. Sample a batch of B episodes, record the states, actions, and rewards
Compute **reward to go**
For each episode, compute:
$$\text{Loss} = -1 * \sum_t [(\text{Log probability of choosing } a_t \text{ at state } S_t) * (\text{Reward_to_go at } t)]$$
2. Average the loss over B episodes.
3. Take a step in the direction of the Gradient of the loss.

