

BASIC ALGORITHMS

Algorithms for small problems

Diego Klabjan
Professor, Industrial Engineering and Management Sciences

Northwestern | McCORMICK SCHOOL OF
ENGINEERING

ENUMERATION

Bellman's Optimality Equation

- $V_t(s_t) = \max_{a_t \in \mathcal{A}} [r(s_t, a_t) + \gamma E_{s_{t+1} \sim p(s_{t+1}|s_t, a_t)} V_{t+1}(s_{t+1}|s_t)]$
- Compute value functions recursively
 - Training (planning)
- Given computed value functions
 - 'Measure' state
 - Solve optimization problem

$$\max_{a \in \mathcal{A}} [r(s, a) + \gamma E_{\bar{s} \sim p(\bar{s}|s, a)} V(\bar{s}|s)]$$

- Often not many actions – enumerate
- Often deterministic system – no expectation

Enumeration

- For $t=T$ down to 0
 - For each possible state s_t
 - Compute

$$V_t(s_t) = \max_{a_t \in \mathcal{A}} [r(s_t, a_t) + \gamma E_{s_{t+1} \sim p(s_{t+1}|s_t, a_t)} V_{t+1}(s_{t+1}|s_t)]$$

- Works if
 - Small number of states
 - Small number of actions
 - Somehow cope with expectation
- Three courses of dimensionality

VALUE ITERATION

Value Iteration

- $V(s) = \max_{a \in \mathcal{A}} [r(s, a) + \gamma E_{\bar{s} \sim p(\bar{s}|s, a)} V(\bar{s}|s)]$
- Assume right-hand side known
 - Use approximate V
 - Can compute left-hand side
 - Gives better approximation of V

Value Iteration

- For $k = 0, 1, 2, \dots$
 - For each possible state s
 - Compute

$$V_{k+1}(s) = \max_{a \in \mathcal{A}} [r(s, a) + \gamma E_{\bar{s} \sim p(\bar{s}|s, a)} V_k(\bar{s}|s)]$$

- If discount factor less than 1 and everything is finite
 - Convergence (pointwise) to optimal value function
- Same pitfalls as enumeration
- No explicit policy

Value Function and Policy

- $\pi(s) = \underset{a \in \mathcal{A}}{\operatorname{argmax}} [r(s, a) + \gamma E_{\bar{s} \sim p(\bar{s}|s, a)} V(\bar{s}|s)]$
 - If V optimal, π is optimal
- $V(s) = r(s, \pi(s)) + \gamma E_{\bar{s} \sim p(\bar{s}|s, \pi(s))} V^\pi(\bar{s}|s)$
 - Must know V^π
 - If π is optimal, V is optimal

Other Applications of Value Iteration

- Shortest Path can be solved by value iteration
- Other algorithms
 - Levensthein distance
 - String algorithms
 - String alignment
 - Dynamic time warping
 - Generalization of Levensthein
 - Graphical models
 - Viterbi algorithm

POLICY ITERATION

Evaluating Policy

- Given policy π find V^π

$$V^\pi(s) = E_{\substack{a \sim \pi(a|s) \\ \bar{s} \sim p(\bar{s}|s,a)}} [r(s, a) + \gamma V^\pi(\bar{s}|s)]$$

- Can use similar idea to value iteration
- Given approximate right-hand side
 - Find better left-hand side by using the equation

Iterative Policy Evaluation

- Problem
 - Evaluate given policy π
- Solution: iterative application of Bellman expectation equation
- $v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_\pi$
- For $k = 0, 1, 2, \dots$
 - For each possible state s compute

$$v_{k+1}(s) = E_{\substack{a \sim \pi(a|s) \\ \bar{s} \sim p(\bar{s}|s,a)}} [r(s, a) + \gamma v_k(\bar{s}|s)]$$

- Convergence to V^π can be proven

Evaluate Policy

$$v_{k+1} = R^\pi + \gamma P^\pi v_k$$

- Limit $k \rightarrow \infty$

$$v^\pi = R^\pi + \gamma P^\pi v^\pi$$

$$v^\pi = (I - \gamma P^\pi)^{-1} R^\pi$$

- Algorithm is a way to compute the inverse

- Inverse exists if discount less than 1

- $R^\pi = \left(E_{a \sim \pi(a|s)} r(s, a) \right)_s = \left(\sum_{a \in \mathcal{A}} \pi(a|s) r(s, a) \right)_s$

- $P^\pi = \left(\sum_{a \in \mathcal{A}} \pi(a|s) P_{ss'}^a \right)_{s,s'}$

Policy Iteration

- Given a policy π
 - Evaluate policy π
$$V^\pi(s) = \mathbb{E}[r_t + \gamma r_{t+1} + \dots | S_t = s]$$
 - Improve the policy by acting greedily with respect to V^π
$$\pi' = \text{greedy}(V^\pi)$$
- This process of policy iteration always converges to π^* - optimal policy
 - Finite cardinality assumptions

Policy Iteration Algorithm

- Loop
 - For $k = 0, 1, 2, \dots$
 - For each possible state s compute

$$v_{k+1}(s) = E_{\substack{a \sim \pi(a|s) \\ \bar{s} \sim p(\bar{s}|s,a)}} [r(s, a) + \gamma v_k(\bar{s}|s)]$$

- Let v^π be the converged function
- For each possible state s compute

$$\pi'(s) = \max_a r(s, a) + \gamma E_{\bar{s} \sim p(\bar{s}|s, a)} v^\pi(\bar{s}|s)$$

- Set $\pi = \pi'$

Generalized Policy Iteration

- Loop
 - For $k = 0, 1, 2, \dots, K$ // K iterations to evaluate the policy
 - For each possible state s compute
$$v_{k+1}(s) = E_{a \sim \pi(a|s)} [r(s, a) + \gamma v_k(\bar{s}|s)]$$
 - Improve the policy by acting greedily with respect to v_{k+1}
$$\pi = \text{greedy}(v_{k+1})$$
- The inner loop approximately computes the inverse of the matrix in
$$\mathbf{v}^\pi = (\mathbf{I} - \gamma \mathbf{P}^\pi)^{-1} \mathbf{R}^\pi$$
- $K=0$
 - Value iteration

- Value iteration

- Per iteration time low
- Needs more iterations

- Policy iteration

- Per iteration time high
 - Controlled by K
- Needs fewer iterations
- More flexible

Weaker convergence assumptions for policy iteration

Trade-off