

POLICY GRADIENT

Peaking popularity
Already the most widely used

Diego Klabjan
Professor, Industrial Engineering and Management Sciences

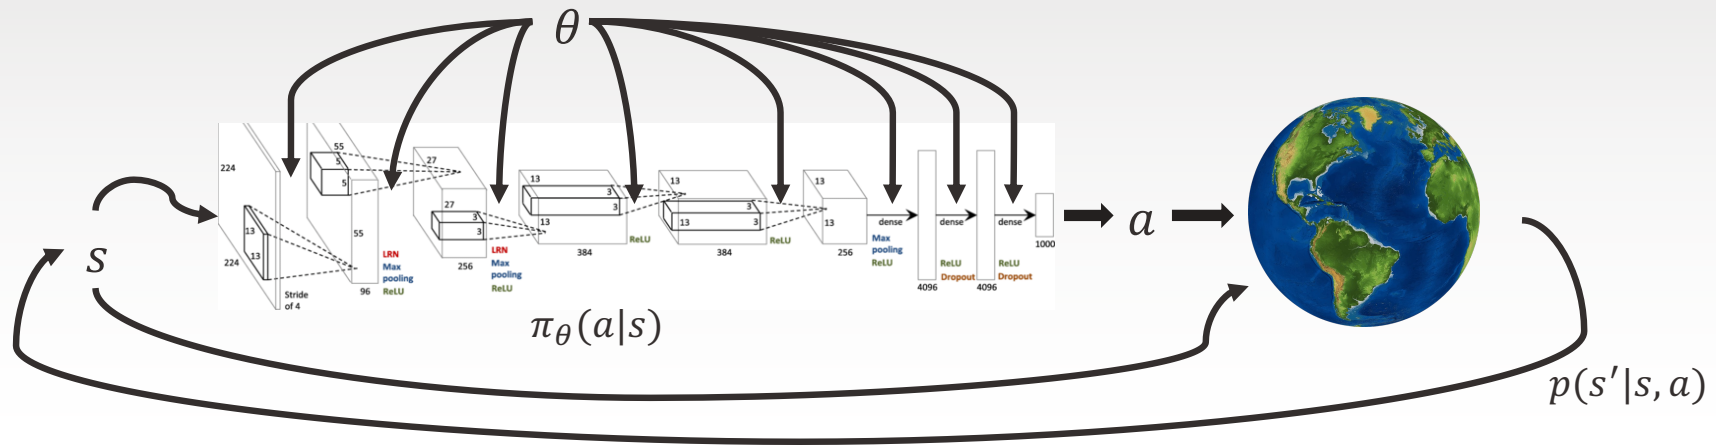
Northwestern | McCORMICK SCHOOL OF
ENGINEERING

Outline

- Formulation
- Algorithm
- Issues
- Bias

FORMULATION

Episodes



$$\underbrace{p_\theta(s_1, a_1, \dots, s_T, a_T)}_{\pi_\theta(\mathcal{T})} = p(s_1) \prod_{t=1}^T \pi_\theta(a_t|s_t) p(s_{t+1}|s_t, a_t)$$

$$\theta^* = \arg \max_{\theta} E_{\mathcal{T} \sim \pi_\theta(\mathcal{T})} \left[\sum_t r(s_t, a_t) \right]$$

Episodes

$$\theta^* = \operatorname{argmax}_{\theta} E_{\tau \sim \pi_{\theta}(\tau)} \left[\sum_t r(s_t, a_t) \right]$$

$$\theta^* = \operatorname{argmax}_{\theta} E_{(s,a) \sim p_{\theta}(s,a)} \left[\sum_t r(s, a) \right]$$

infinite horizon case

$$\theta^* = \operatorname{argmax}_{\theta} \sum_{t=1}^T E_{(s_t, a_t) \sim p_{\theta}(s_t, a_t)} [r(s_t, a_t)]$$

finite horizon case

Telescoping

$$\theta^* = \operatorname{argmax}_{\theta} \underbrace{E_{\tau \sim \pi_{\theta}(\tau)} \left[\sum_t r(s_t, a_t) \right]}_{J(\theta)}$$

- Assuming everything is finite
 - Otherwise an approximation

$$J(\theta) = E_{\tau \sim \pi_{\theta}(\tau)} \left[\sum_t r(s_t, a_t) \right] = \sum_i p(\tau^i) \sum_t r(s_t^i, a_t^i)$$

sum over samples from π_{θ}

ALGORITHM

Gradient Optimization

- $J(\theta)$ not convex
- Gradient optimization still applicable
$$\theta = \theta + \alpha \nabla J(\theta)$$
 - α learning rate
- Challenge to find the gradient
- $g(\theta) = E_{z \sim Q} f(\theta, z)$
 - $\nabla g(\theta) = E_{z \sim Q} \nabla_{\theta} f(\theta, z)$
- What about $g(\theta) = E_{z \sim Q(\theta)} f(\theta, z)$?

Direct Policy Differentiation

$$\theta^* = \operatorname{argmax}_{\theta} \underbrace{E_{\tau \sim \pi_{\theta}(\tau)} \left[\sum_t r(s_t, a_t) \right]}_{J(\theta)}$$

a convenient identity

$$\pi_{\theta}(\tau) \nabla_{\theta} \log \pi_{\theta}(\tau) = \pi_{\theta}(\tau) \frac{\nabla_{\theta} \pi_{\theta}(\tau)}{\pi_{\theta}(\tau)} = \nabla_{\theta} \pi_{\theta}(\tau)$$

$$J(\theta) = E_{\tau \sim \pi_{\theta}(\tau)} \underbrace{[r(\tau)]}_{\sum_{t=1}^T r(s_t, a_t)} = \int \pi_{\theta}(\tau) r(\tau) d\tau$$

$$\nabla_{\theta} J(\theta) = \int \nabla_{\theta} \pi_{\theta}(\tau) r(\tau) d\tau = \int \pi_{\theta}(\tau) \nabla_{\theta} \log \pi_{\theta}(\tau) r(\tau) d\tau = E_{\tau \sim \pi_{\theta}(\tau)} [\nabla_{\theta} \log \pi_{\theta}(\tau) r(\tau)]$$

Direct Policy Differentiation

$$\theta^* = \underset{\theta}{\operatorname{argmax}} J(\theta)$$

$$J(\theta) = E_{\tau \sim \pi_{\theta}(\tau)}[r(\tau)]$$

$$\nabla_{\theta} J(\theta) = E_{\tau \sim \pi_{\theta}(\tau)}[\nabla_{\theta} \log \pi_{\theta}(\tau) r(\tau)]$$

$$\nabla_{\theta} \left[\cancel{\log p(s_1)} + \sum_{t=1}^T \log \pi_{\theta}(a_t | s_t) + \cancel{\log p(s_{t+1} | s_t, a_t)} \right]$$

log of both sides

$$\underbrace{p_{\theta}(s_1, a_1, \dots, s_T, a_T)}_{\pi_{\theta}(\tau)} = p(s_1) \prod_{t=1}^T \pi_{\theta}(a_t | s_t) p(s_{t+1} | s_t, a_t)$$

$$\log \pi_{\theta}(\tau) = \log p(s_1) + \sum_{t=1}^T \log \pi_{\theta}(a_t | s_t) + \log p(s_{t+1} | s_t, a_t)$$

$$\nabla_{\theta} J(\theta) = E_{\tau \sim \pi_{\theta}(\tau)} \left[\left(\sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right) \left(\sum_{t=1}^T r(s_t, a_t) \right) \right]$$

Algorithm

$$J(\theta) = E_{\tau \sim p_{\theta}(\tau)} \left[\sum_t r(s_t, a_t) \right] \approx \frac{1}{N} \sum_i \sum_t r(s_t^i, a_t^i)$$

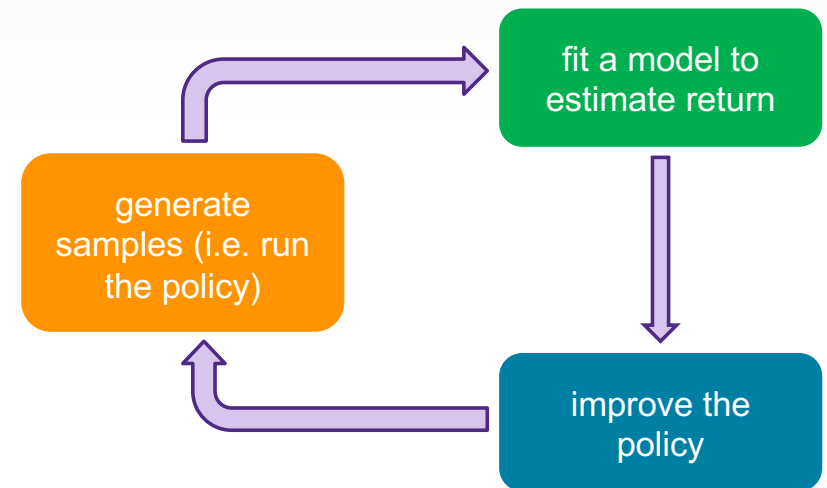
$$\nabla_{\theta} J(\theta) = E_{\tau \sim \pi_{\theta}(\tau)} \left[\left(\sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right) \left(\sum_{t=1}^T r(s_t, a_t) \right) \right]$$

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \left(\sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_t^i | s_t^i) \right) \left(\sum_{t=1}^T r(s_t^i, a_t^i) \right)$$

$$\theta \leftarrow \theta + \alpha \nabla_{\theta} J(\theta)$$

REINFORCE algorithm:


1. sample $\{\tau^i\}$ from $\pi_{\theta}(a_t | s_t)$ (run the policy)
2. $\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_i \left(\sum_t \nabla_{\theta} \log \pi_{\theta}(a_t^i | s_t^i) \right) \left(\sum_t r(s_t^i | a_t^i) \right)$
3. $\theta \leftarrow \theta + \alpha \nabla_{\theta} J(\theta)$



Algorithm

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \left(\sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_t^i | s_t^i) \right) \left(\sum_{t=1}^T r(s_t^i, a_t^i) \right)$$

what is this?



- Network produces softmax
- Take log as “loss” function
 - Standard back propagation

Multi-step MDP



- Replace instantaneous reward r with long-term value $Q^\pi(s, a)$
- For discounted objective
 - Any differentiable policy $\pi_\theta(a|s)$
 - Can be shown

$$\nabla_\theta J(\theta) = \mathbb{E}_{\pi_\theta} [\nabla_\theta \log \pi_\theta(a|s) Q^{\pi_\theta}(s, a)]$$


ENHANCEMENTS

Variance of Policy-Gradient

- Known to have high variance



Created by Ghani Ebrahimi



Baselines

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \nabla_{\theta} \log \pi_{\theta}(\tau^i) [r(\tau^i) - b]$$

a convenient identity

$$\pi_{\theta}(\tau) \nabla_{\theta} \log \pi_{\theta}(\tau) = \nabla_{\theta} \pi_{\theta}(\tau)$$

- Let b (baseline) be constant

- Independent of τ

$$E[\nabla_{\theta} \log \pi_{\theta}(\tau) b] = \int \pi_{\theta}(\tau) \nabla_{\theta} \log \pi_{\theta}(\tau) b \, d\tau = \int \nabla_{\theta} \pi_{\theta}(\tau) b \, d\tau = b \nabla_{\theta} \int \pi_{\theta}(\tau) \, d\tau = b \nabla_{\theta} 1 = 0$$

- Subtracting a baseline is unbiased in expectation
- Candidate

$$b = \frac{1}{N} \sum_{i=1}^N r(\tau^i)$$

Reducing Variance

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \left(\sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_t^i | s_t^i) \right) \left(\sum_{t=1}^T r(s_t^i, a_t^i) \right)$$

- Policy at time t' cannot affect reward at time t when $t < t'$

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_t^i | s_t^i) \underbrace{\left(\sum_{t'=1}^T r(s_{t'}^i, a_{t'}^i) \right)}_{\substack{\text{"reward to go"} \\ \hat{Q}_{i,0}}}$$

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_t^i | s_t^i) \left(\sum_{\substack{t'=t \\ \text{red}}}^T r(s_{t'}^i, a_{t'}^i) \right) \text{ REINFORCE algorithm}$$

Off-policy PG and Importance Sampling

- Recall $\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(a|s) Q^{\pi_{\theta}}(s, a)]$
 - At time t
 - Q matters
 - Q captures reward only from t onwards
- Scientific justification for truncation

Analyzing Variance

$$\text{Var}[x] = E[x^2] - E[x]^2$$

$$\nabla_{\theta} J(\theta) = E_{\tau \sim \pi_{\theta}(\tau)} [\nabla_{\theta} \log \pi_{\theta}(\tau) (r(\tau) - b)]$$

$$\begin{aligned} \text{Var} &= E_{\tau \sim \pi_{\theta}(\tau)} [(\nabla_{\theta} \log \pi_{\theta}(\tau) (r(\tau) - b))^2] - \underbrace{E_{\tau \sim \pi_{\theta}(\tau)} [\nabla_{\theta} \log \pi_{\theta}(\tau) (r(\tau) - b)]^2}_{[E_{\tau \sim \pi_{\theta}(\tau)} [\nabla_{\theta} \log \pi_{\theta}(\tau) r(\tau)]]^2} \end{aligned}$$

$$\begin{aligned} \frac{d\text{Var}}{db} &= \frac{d}{db} E[g(\tau)^2 (r(\tau) - b)^2] = \frac{d}{db} (E[g(\tau)^2 r(\tau)^2] - 2E[g(\tau)^2 r(\tau) b] + b^2 E[g(\tau)^2]) \\ &= -2E[g(\tau)^2 r(\tau)] + 2bE[g(\tau)^2] = 0 \end{aligned}$$

$$b = \frac{E[g(\tau)^2 r(\tau)]}{E[g(\tau)^2]}$$

← This is just expected reward, but weighted by gradient magnitudes!

Optimal constant baseline

Reducing Variance Using a Baseline

- Good baseline
 - State value function $B(s) = V^{\pi_\theta}(s)$
- Policy gradient

$$A^{\pi_\theta}(s, a) = Q^{\pi_\theta}(s, a) - V^{\pi_\theta}(s)$$

$$\nabla_\theta J(\theta) = \mathbb{E}_{\pi_\theta} [\nabla_\theta \log \pi_\theta(s, a) A^{\pi_\theta}(s, a)]$$

- Advantage function $A^{\pi_\theta}(s, a)$
 - For optimal policy advantage function always non-positive

Estimating Advantage Function

- Two function approximators and two parameter vectors

$$\begin{aligned}V_v(s) &\approx V^{\pi_\theta}(s) \\ Q_w(s, a) &\approx Q^{\pi_\theta}(s, a) \\ A_{w,v}(s, a) &= Q_w(s, a) - V_v(s)\end{aligned}$$

- θ fixed; w, v learnable
 - Updating both value functions by ideas from Q-learning
- Alternative to have a single network for advantage

Policy Gradient Algorithms

- Many equivalent forms

$$\begin{aligned}\nabla_{\theta} J(\theta) &= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) \textcolor{red}{r}_t] && \text{REINFORCE} \\ &= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) \textcolor{red}{Q}^w(s, a)] && \text{Q Actor-Critic} \\ &= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) \textcolor{red}{A}^{w,v}(s, a)] && \text{Advantage Actor-Critic}\end{aligned}$$

- Equality for expectation
- When sampling and approximating functions
 - Different algorithms

Algorithm

- Loop
 1. sample $\{\tau^i\}$ from $\pi_\theta(a_t|s_t)$ (run the policy)
 2. $\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_i (\sum_t \nabla_\theta \log \pi_\theta(a_t^i|s_t^i)) (\sum_t A_{w,v}(s_t^i, a_t^i))$
 3. $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$
 4. Create set: episode τ^i and each (s_t^i, a_t^i)
 5. Adjust w, v based on MSE (recomputing the advantage function)
- Only value and Q functions approximation used
 - MSE based only on states and policy (actions)

UPDATING FUNCTIONS

Functional Approximation of V

- Updating the value function approximation
- At state s
 - Have reward plus future based on approximate value function
 - Based on optimal action
 - Have value of approximate value function
- Match them
 - L2 loss

fitted value iteration algorithm:

1. set $y_i \leftarrow \max_{a_i}(r(s_i, a_i) + \gamma E[V_v(s'_i)])$
2. set $v \leftarrow \operatorname{argmin}_v, \frac{1}{2} \sum_i (V_v(s_i) - y_i)^2$

Functional Approximation of Q


fitted value iteration algorithm:

- 1. set $y_i^1 \leftarrow \max_{a_i} (r(s_i, a_i) + \gamma E[V_v(s'_i)])$
- 2. set $v \leftarrow \operatorname{argmin}_v, \frac{1}{2} \sum_i (V_v(s_i) - y_i^1)^2$

fitted Q iteration algorithm:

- 1. set $y_i^2 \leftarrow r(s_i, a_i) + \gamma E[V_v(s'_i)] \longleftarrow \text{approximate } E[V_v(s'_i)] \approx \max_{a'} Q_w(s'_i, a'_i)$
- 2. set $w \leftarrow \operatorname{argmin}_w, \frac{1}{2} \sum_i (Q_w(s_i, a_i) - y_i^2)^2$

Doesn't require simulation of actions!

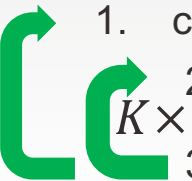


Value Iteration with Fitted Q-factor


- Observed data are trajectories
 - Formally, $U = \{(s_i, a_i, s'_i, r_i) | i \in N\}$
- Loop
 - Sample $S \subseteq U$
 - For $i \in S$
 - $y_i \leftarrow r_i + \gamma \max_{a'_i} Q_w(s'_i, a'_i)$
 - Set $w \leftarrow \operatorname{argmin}_w, \frac{1}{2} \sum_{i \in S} (Q_w(s_i, a_i) - y_i)^2$

Online Version

full fitted Q-iteration algorithm:

- 
1. collect dataset $\{(s_i, a_i, s'_i, r_i)\}$ using some policy
 2. set $y_i \leftarrow r(s_i, a_i) + \gamma \max_{a'_i} Q_w(s'_i, a'_i)$
 3. set $w \leftarrow \operatorname{argmin}_w, \frac{1}{2} \sum_i (Q_w(s_i, a_i) - y_i)^2$

online Q iteration algorithm:

- 
1. take some action a_i and observe (s_i, a_i, s'_i, r_i)
 2. $y_i = r(s_i, a_i) + \gamma \max_{a'_i} Q_w(s'_i, a'_i)$
 3. $w \leftarrow w - \alpha \frac{dQ_w}{dw}(s_i, a_i)(Q_w(s_i, a_i) - y_i)$

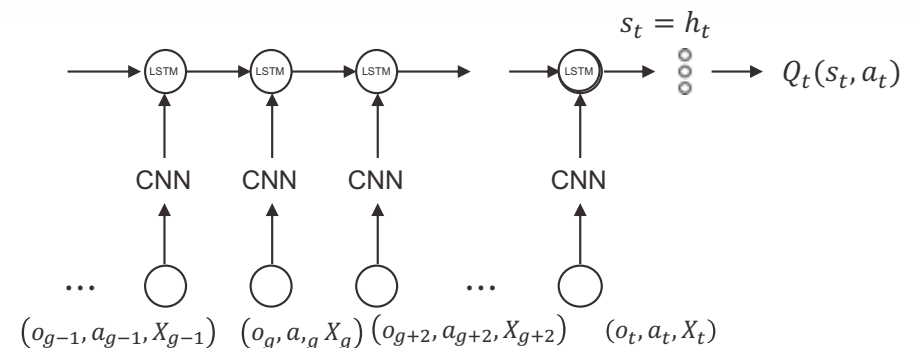

Overall Algorithm

- Loop

1. sample $\{\tau^i\}$ from $\pi_\theta(a_t|s_t)$ (run the policy)
2. $\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_i (\sum_t \nabla_\theta \log \pi_\theta(a_t^i|s_t^i)) (\sum_t A_{w,v}(s_t^i, a_t^i))$
3. $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$
4. Create set: episode τ^i and each (s_t^i, a_t^i)
5. For all t, i present in episodes in step 4
 1. Let $y_{ti}^1 \leftarrow \max_{a_i} (r(s_t^i, a_t^i) + \gamma E[V_v(s_{t'}^i)])$.
 2. Let $y_{ti}^2 \leftarrow r(s_t^i, a_t^i) + \gamma E[V_v(s_{t'}^i)]$. // OR $y_{ti}^2 \leftarrow r(s_t^i, a_t^i) + \gamma \max_{a_i'} Q_w(s_{t'}^i, a_{t'}^i)$
6. Value function update: set $v \leftarrow \operatorname{argmin}_v, \frac{1}{2} \sum_{ti} (V_v(s_t^i) - y_{ti}^1)^2$
7. Q-factor update: set $w \leftarrow \operatorname{argmin}_w, \frac{1}{2} \sum_{ti} (Q_w(s_t^i, a_t^i) - y_{ti}^2)^2$

Deep Q-Learning

- Deep network used to model Q factor
- State modeled by RNN
 - Possibly combined with CNN if images involved
- Final network a combination of CNN+RNN+Fully connected



Exploration-Exploitation

- Exploration-Exploitation tradeoff
- Have visited part of the state space and found a reward of 100
 - Is this the best we can hope for?
 - Should we keep 'pounding' the visited states and figure out actions that lead to 100?
 - Perhaps there are other states that we have not yet visited that lead to even higher reward
- Exploitation
 - Should we stick with what we know
 - Find a good policy with respect to this knowledge
 - Risk of missing out on a better reward somewhere else
- Exploration
 - Should we look for states with more reward
 - At risk of wasting time and getting some negative reward

Exploration-Exploitation Epsilon-Greedy

- Exploitation
 - Follow your current best policy
 - Be greedy
- Exploration
 - Deviate to explore
 - Strategy
 - Greedy policy vector a
 - Perturb it by epsilon
 $a + \epsilon$
 - ϵ random vector
 - Has to be probability vector
- Common exploration
 - Probability(a)
 - $1 - \epsilon$ if $a = \text{optimal (argmax)}$
 - $\epsilon / (|\mathcal{A}| - 1)$ otherwise

Exploration-Exploitation

- Sample from one policy
 - Exploration
 - Behavior policy
- Update and optimization different policy
 - Exploitation
 - Learning policy
- Same policies: on-policy learning
- Two different policies: off-policy learning

Off-policy vs On-policy

- On-policy algorithm
 - Learn the policy being executed by the agent
 - One policy
- Off-policy algorithm
 - Evaluate a policy from samples generated by a different policy
 - Target policy
 - Learn policy independent of policy taken by agent (behavior policy)
 - Two policies

Off-policy Framework

- $\pi_{\theta'}$ is target policy
 - Gradient with respect to this policy
 - Gradient at this point
- Loop
 - Sample episode based on policy π_{θ}
 - // Behavior policy
 - Perform gradient step at $\pi_{\theta'}$
 - // Adjusts $\pi_{\theta'}$
- Challenge
 - The two policies should be somehow related
 - Idea: gradient taken with respect to a function that involves π_{θ}

Policy Gradient and On-policy

$$\theta^* = \operatorname{argmax}_{\theta} J(\theta)$$

$$J(\theta) = E_{\tau \sim \pi_{\theta}(\tau)}[r(\tau)]$$

$$\nabla_{\theta} J(\theta) = E_{\tau \sim \pi_{\theta}(\tau)}[\nabla_{\theta} \log \pi_{\theta}(\tau) r(\tau)]$$

Same policy

- Neural networks change only slightly with each gradient step
- On-policy learning inefficient

REINFORCE algorithm:

1. sample $\{\tau^i\}$ from $\pi_{\theta}(a_t|s_t)$
2. $\nabla_{\theta} J(\theta) \approx \sum_i (\sum_t \nabla_{\theta} \log \pi_{\theta}(a_t^i|s_t^i)) (\sum_t r(s_t^i, a_t^i))$
3. $\theta \leftarrow \theta + \alpha \nabla_{\theta} J(\theta)$

Must have this

Off-policy PG and Importance Sampling

$$\theta^* = \operatorname{argmax}_{\theta} J(\theta)$$

$$J(\theta') = E_{\tau \sim \pi_{\theta'}(\tau)}[r(\tau)]$$

We have samples from some $\pi_{\theta}(\tau)$

$$J(\theta') = E_{\tau \sim \pi_{\theta}(\tau)} \left[\frac{\pi_{\theta'}(\tau)}{\pi_{\theta}(\tau)} r(\tau) \right]$$

$$\pi_{\theta}(\tau) = p(s_1) \prod_{t=1}^T \pi_{\theta}(a_t | s_t) p(s_{t+1} | s_t, a_t)$$

$$\frac{\pi_{\theta'}(\tau)}{\pi_{\theta}(\tau)} = \frac{\cancel{p(s_1)} \prod_{t=1}^T \pi_{\theta'}(a_t | s_t) \cancel{p(s_{t+1} | s_t, a_t)}}{\cancel{p(s_1)} \prod_{t=1}^T \pi_{\theta}(a_t | s_t) \cancel{p(s_{t+1} | s_t, a_t)}} = \frac{\prod_{t=1}^T \pi_{\theta'}(a_t | s_t)}{\prod_{t=1}^T \pi_{\theta}(a_t | s_t)}$$

Importance sampling

$$\begin{aligned} E_{x \sim p(x)}[f(x)] &= \int p(x) f(x) dx \\ &= \int \frac{q(x)}{q(x)} p(x) f(x) dx \\ &= \int q(x) \frac{p(x)}{q(x)} f(x) dx \\ &= E_{x \sim q(x)} \left[\frac{p(x)}{q(x)} f(x) \right] \end{aligned}$$

Off-policy PG and Importance Sampling

$$\theta^* = \operatorname{argmax}_{\theta} J(\theta)$$

$$J(\theta) = E_{\tau \sim \pi_{\theta}(\tau)}[r(\tau)]$$

Convenient identity

$$\pi_{\theta}(\tau) \nabla_{\theta} \log \pi_{\theta}(\tau) = \nabla_{\theta} \pi_{\theta}(\tau)$$

Estimate the value of some *new* parameters θ' :

$$J(\theta') = E_{\tau \sim \pi_{\theta}(\tau)} \left[\frac{\pi_{\theta'}(\tau)}{\pi_{\theta}(\tau)} r(\tau) \right]$$

Depends on θ'

$$\nabla_{\theta'} J(\theta') = E_{\tau \sim \pi_{\theta}(\tau)} \left[\frac{\nabla_{\theta'} \pi_{\theta'}(\tau)}{\pi_{\theta}(\tau)} r(\tau) \right] = E_{\tau \sim \pi_{\theta}(\tau)} \left[\frac{\pi_{\theta'}(\tau)}{\pi_{\theta}(\tau)} \nabla_{\theta'} \log \pi_{\theta'}(\tau) r(\tau) \right]$$

At $\theta = \theta'$

$$\nabla_{\theta} J(\theta) = E_{\tau \sim \pi_{\theta}(\tau)} [\nabla_{\theta} \log \pi_{\theta}(\tau) r(\tau)] - \text{Original policy gradient}$$

Off-policy PG and Importance Sampling

$$\theta^* = \underset{\theta}{\operatorname{argmax}} J(\theta)$$

$$J(\theta) = E_{\tau \sim \pi_{\theta}(\tau)}[r(\tau)]$$

$$\frac{\pi_{\theta'}(\tau)}{\pi_{\theta}(\tau)} = \frac{\prod_{t=1}^T \pi_{\theta'}(a_t | s_t)}{\prod_{t=1}^T \pi_{\theta}(a_t | s_t)}$$

$$\nabla_{\theta'} J(\theta') = E_{\tau \sim \pi_{\theta}(\tau)} \left[\frac{\pi_{\theta'}(\tau)}{\pi_{\theta}(\tau)} \nabla_{\theta'} \log \pi_{\theta'}(\tau) r(\tau) \right]$$

$$= E_{\tau \sim \pi_{\theta}(\tau)} \left[\left(\prod_{t=1}^T \frac{\pi_{\theta'}(a_t | s_t)}{\pi_{\theta}(a_t | s_t)} \right) \left(\sum_{t=1}^T \nabla_{\theta'} \log \pi_{\theta'}(a_t | s_t) \right) \left(\sum_{t=1}^T r(s_t, a_t) \right) \right]$$

$$= E_{\tau \sim \pi_{\theta}(\tau)} \left[\sum_{t=1}^T \nabla_{\theta'} \log \pi_{\theta'}(a_t | s_t) \left(\prod_{t'=1}^t \frac{\pi_{\theta'}(a_{t'} | s_{t'})}{\pi_{\theta}(a_{t'} | s_{t'})} \right) \left(\sum_{t'=t}^T r(s_{t'}, a_{t'}) \right) \right]$$

future has no
effect on present

at t past reward
does not matter
(sunk cost)

Off-policy PG and Importance Sampling

$$\nabla_{\theta'} J(\theta') = E_{\tau \sim \pi_{\theta}(\tau)} \left[\sum_{t=1}^T \nabla_{\theta'} \log \pi_{\theta'}(a_t | s_t) \underbrace{\left(\prod_{t'=1}^t \frac{\pi_{\theta'}(a_{t'} | s_{t'})}{\pi_{\theta}(a_{t'} | s_{t'})} \right)}_{\text{exponential in } T \dots} \left(\sum_{t'=t}^T r(s_{t'}, a_{t'}) \right) \right]$$

$$\theta^* = \operatorname{argmax}_{\theta} \sum_{t=1}^T E_{(s_t, a_t) \sim p_{\theta}(s_t, a_t)} [r(s_t, a_t)]$$

$$J(\theta) = \sum_{t=1}^T E_{(s_t, a_t) \sim p_{\theta}(s_t, a_t)} [r(s_t, a_t)] = \sum_{t=1}^T E_{s_t \sim p_{\theta}(s_t)} \left[E_{a_t \sim \pi_{\theta}(a_t | s_t)} [r(s_t, a_t)] \right]$$

$$J(\theta') = \sum_{t=1}^T E_{s_t \sim p_{\theta}(s_t)} \left[\cancel{\frac{p_{\theta'}(s_t)}{p_{\theta}(s_t)}} E_{a_t \sim \pi_{\theta}(a_t | s_t)} \left[\frac{\pi_{\theta'}(a_t | s_t)}{\pi_{\theta}(a_t | s_t)} r(s_t, a_t) \right] \right]$$

Ignore this part

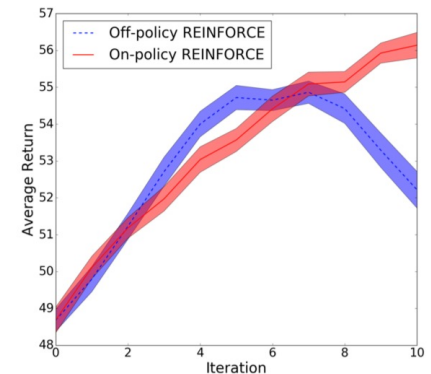
Perform
importance
sampling for
states and
actions

Off-policy PG Algorithm

Off-policy REINFORCE algorithm:

1. sample $\{\tau^i\}$ from $\pi_\theta(a_t|s_t)$ // behavior policy
2. $\nabla_{\theta'} J(\theta') \approx \sum_i \left(\sum_t \nabla_{\theta'} \log \pi_{\theta'}(a_t^i | s_t^i) r(s_t^i, a_t^i) / \pi_\theta(a_t^i | s_t^i) \right)$
3. $\theta' \leftarrow \theta' + \alpha \nabla_{\theta'} J(\theta')$ // target policy

- Choice of behavior policy
 - Epsilon greedy with respect to learned policy
 - Policy that minimizes the variance of learned policy
 - Details complicated




Hanna and Stone 2018: Towards a Data Efficient Off-Policy Policy Gradient

Practical Version

- Replace reward with advantage

Off-policy actor-critic algorithm:

- 
1. sample $\{\tau^i\}$ from $\pi_\theta(a_t|s_t)$ // behavior policy
 2. $\nabla_{\theta'} J(\theta') \approx \sum_i \left(\sum_t \nabla_{\theta'} \log \pi_{\theta'}(a_t^i | s_t^i) A^{\pi_\theta}(s_t^i, a_t^i) / \pi_\theta(a_t^i | s_t^i) \right)$
 3. $\theta' \leftarrow \theta' + \alpha \nabla_{\theta'} J(\theta')$ // target policy

Policy Gradient in Practice

- Remember that the gradient has high variance
 - Not the same as supervised learning
 - Much more uncertainty
 - Gradients will be really noisy
- Consider using much larger batches
- Tweaking learning rates is very hard
 - Adaptive step size rules like ADAM help

Advantages of Policy-Based RL

- Advantages:
 - Good convergence properties
 - Effective in high-dimensional or continuous action spaces
 - Learns stochastic policies
- Disadvantages:
 - Typically converge to a local rather than global optimum
 - Evaluating a policy is typically inefficient and high variance